

Sprawozdanie

Piotr Bogdanowicz

Section #1

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
cd4463ae4b5f	bridge	bridge	local
afdf9c132cee	host	host	local
05f36a05cc57	none	null	local

```
$ docker network inspect bridge
```

```
[
  {
    "Name": "bridge",
    "Id": "cd4463ae4b5f786dbcff4990eb048911632b7d04b1e79dbe70b5d6384bc83d56",
    "Created": "2018-12-03T19:46:42.752299021Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.17.0.0/16"
        }
      ]
    }
  }
]
```

```

    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {
      "com.docker.network.bridge.default_bridge": "true",
      "com.docker.network.bridge.enable_icc": "true",
      "com.docker.network.bridge.enable_ip_masquerade": "true",
      "com.docker.network.bridge.host_binding_ipv4": "0.0.0.0",
      "com.docker.network.bridge.name": "docker0",
      "com.docker.network.driver.mtu": "1500"
    },
    },
    "Labels": {}
  }
}

```

```

$ docker info
Containers: 0
Running: 0
Paused: 0
Stopped: 0
Images: 0
Server Version: 18.06.1-ce
Storage Driver: overlay2
Backing Filesystem: xfs
Supports d_type: true
Native Overlay Diff: true
Logging Driver: json-file
Cgroup Driver: cgroupfs
Plugins:
Volume: local
Network: bridge host ipvlan macvlan null overlay
Log: awslogs fluentd gcplogs gelf journald json-file logentries splu
nk syslog

```

Section #2

```

$ docker network ls

```

NETWORK ID	NAME	DRIVER	SCOPE
cd4463ae4b5f	bridge	bridge	local
afdf9c132cee	host	host	local
05f36a05cc57	none	null	local

```
$ apk add bridge
(1/1) Installing bridge (1.5-r3)
OK: 302 MiB in 111 packages
[node1] (local) root@192.168.0.27 ~
$ brctl show
```

bridge name	bridge id	STP enabled	interfaces
docker0	8000.0242aa4e8eea	no	

```
2: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN
    link/ether 02:42:aa:4e:8e:ea brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
```

```
$ docker run -dt ubuntu sleep infinity
Unable to find image 'ubuntu:latest' locally
latest: Pulling from library/ubuntu
32802c0cfa4d: Pull complete
da1315cffa03: Pull complete
fa83472a3562: Pull complete
f85999a86bef: Pull complete
Digest: sha256:6d0e0c26489e33f5a6f0020edface2727db9489744ecc9b4f50c7fa671f23c49
Status: Downloaded newer image for ubuntu:latest
09c19697e6c5660ddbdc2aa749a3fa04d432ca3c15a48b1e97570ad47395a683
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
09c19697e6c5	ubuntu	"sleep infinity"	About a minute ago
Up About a minute		frosty_kilby	

```
$ brctl show
```

bridge name	bridge id	STP enabled	interfaces
docker0	8000.0242aa4e8eea	no	veth2a98424

For docker0:

```
"Containers": {  
  "09c19697e6c5660ddbdc2aa749a3fa04d432ca3c15a48b1e97570ad47395a683"  
": {  
  "Name": "frosty_kilby",  
  "EndpointID": "f6215adba0ed92898f0d0fca27dad052c92652a79cb5b1  
f13f4ed4126e74244d",  
  "MacAddress": "02:42:ac:11:00:02",  
  "IPv4Address": "172.17.0.2/16",  
  "IPv6Address": ""  
  }  
},
```

```
$ ping -c5 172.17.0.2  
PING 172.17.0.2 (172.17.0.2): 56 data bytes  
64 bytes from 172.17.0.2: seq=0 ttl=64 time=0.115 ms  
64 bytes from 172.17.0.2: seq=1 ttl=64 time=0.059 ms  
64 bytes from 172.17.0.2: seq=2 ttl=64 time=0.079 ms  
64 bytes from 172.17.0.2: seq=3 ttl=64 time=0.063 ms  
64 bytes from 172.17.0.2: seq=4 ttl=64 time=0.096 ms  
  
--- 172.17.0.2 ping statistics ---  
5 packets transmitted, 5 packets received, 0% packet loss  
round-trip min/avg/max = 0.059/0.082/0.115 ms
```

```
root@09c19697e6c5:/#
```

```
$ docker exec -it 09c19697e6c5 /bin/bash  
root@09c19697e6c5:/# apt-get update && apt-get install -y iputi  
ls-ping  
Get:1 http://security.ubuntu.com/ubuntu bionic-security InRelea  
se [83.2 kB]  
Get:2 http://archive.ubuntu.com/ubuntu bionic InRelease [242 kB  
]
```

```

root@09c19697e6c5:/# ping -c5 www.github.com
PING github.com (192.30.253.112) 56(84) bytes of data.
64 bytes from lb-192-30-253-112-iad.github.com (192.30.253.112): icmp_seq=1 ttl=50 time=1.90 ms
64 bytes from lb-192-30-253-112-iad.github.com (192.30.253.112): icmp_seq=2 ttl=50 time=1.77 ms
64 bytes from lb-192-30-253-112-iad.github.com (192.30.253.112): icmp_seq=3 ttl=50 time=1.87 ms
64 bytes from lb-192-30-253-112-iad.github.com (192.30.253.112): icmp_seq=4 ttl=50 time=1.89 ms
64 bytes from lb-192-30-253-112-iad.github.com (192.30.253.112): icmp_seq=5 ttl=50 time=1.87 ms

--- github.com ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 1.772/1.863/1.906/0.067 ms
root@09c19697e6c5:/#

```

```

$ docker stop 09c19697e6c5
09c19697e6c5
[node1] (local) root@192.168.0.27 ~
$ docker ps

```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

```

$ docker run --name web1 -d -p 8080:80 nginx
Unable to find image 'nginx:latest' locally
latest: Pulling from library/nginx
a5a6f2f73cd8: Pull complete
1ba02017c4b2: Pull complete
33b176c904de: Pull complete
Digest: sha256:5d32f60db294b5deb55d078cd4feb410ad88e6fe77500c87d3970eca97f54dba
Status: Downloaded newer image for nginx:latest
860d8bacf95b8264feaae4e67eb0323eeace1e29bae6f048a063d409a8159cb
0

```

```

$ docker ps

```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
860d8bacf95b	nginx	"nginx -g 'daemon of...'"	19 seconds ago	Up 18 seconds	0.0.0.0:8080->80/tcp	web1

```
$ curl 127.0.0.1:8080
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully i
nstalled and
```

Section #3

Section #3 - Overlay Networking

Step 1: The Basics

In this step you'll initialize a new Swarm, join a single worker node, and verify the operations worked.

Run `docker swarm init --advertise-addr $(hostname -i)`.

```
docker swarm init --advertise-addr $(hostname -i)
```

```
Swarm initialized: current node (rzyy572arjko2w0j82zvjk6u) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join \
--token SWMTKN-1-69b2x1u2wtjdmot0oqxjw1r2d27f0lbnhbdhy83chln1f6es5-37ykdpu0vylenefe2439cpgf \
10.0.0.5:2377
```

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

In the first terminal copy the entire `docker swarm join ...` command that is displayed as part of the output from your terminal output. Then, paste the copied command into the second terminal.

```
docker swarm join \
> --token SWMTKN-1-69b2x1u2wtjdmot0oqxjw1r2d27f0lbnhbdhy83chln1f6es5-37ykdpu0vylenefe2439cpgf \
> 10.0.0.5:2377
This node joined a swarm as a worker.
```

Run a `docker node ls` to verify that both nodes are part of the Swarm.

```
docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
ijmqghkdyaf5h9rjzyngdm48	node2	Ready	Active	
rzyy572arjko2w0j82zvjk6u	node1	Ready	Active	Leader

The `ID` and `HOSTNAME` values may be different in your lab. The important thing to check is that both nodes have joined the Swarm and are ready and active.

Step 2: Create an overlay network

Step 2: Create an overlay network

Now that you have a Swarm initialized it's time to create an **overlay** network.

Create a new overlay network called "overnet" by running `docker network create -d overlay overnet`.

```
docker network create -d overlay overnet
```

```
wlqmvajmzskn84bqbd1yuy
```

Use the `docker network ls` command to verify the network was created successfully.

```
docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
3430a6920bf	bridge	bridge	local
a4d58435009	docker_gwbridge	bridge	local
a7449465c379	host	host	local
8hq1n8nak54x	ingress	overlay	swarm
06c349b9cc77	none	null	local
wlqmvajmzsk	overnet	overlay	swarm

The new "overnet" network is shown on the last line of the output above. Notice how it is associated with the **overlay** driver and is scoped to the entire Swarm.

NOTE: The other new networks (ingress and docker_gwbridge) were created automatically when the Swarm cluster was created.

Run the same `docker network ls` command from the second terminal.

```
docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
55f10b3b0ed	bridge	bridge	local
b7b30433a639	docker_gwbridge	bridge	local
a7449465c379	host	host	local
8hq1n8nak54x	ingress	overlay	swarm
06c349b9cc77	none	null	local

Notice that the "overnet" network does **not** appear in the list. This is because Docker only extends overlay networks to hosts when they are needed. This is usually when a host runs a task from a service that is created on the network. We will see this shortly.

```
if the commandline doesn't appear in the terminal, make sure popups are enabled or try resizing
```

```
ID          HOSTNAME          STA
TUS        AVAILABILITY     MANAGER STATUS
ENGINE VERSION
tg9lkohjqwwtadv0208vzsv2 * node1      Rea
dy          Active          Leader
18.06.1-ce
(node1) (local) root@192.168.0.22 ~
$ docker node ls
ID          HOSTNAME          STA
TUS        AVAILABILITY     MANAGER STATUS
ENGINE VERSION
tg9lkohjqwwtadv0208vzsv2 * node1      Rea
dy          Active          Leader
18.06.1-ce
q2m74vmo7u9z6n9ovq407thru node2      Rea
dy          Active
18.06.1-ce
(node1) (local) root@192.168.0.22 ~
$
w04jrjfoxm211nxuwiomye5wgl7fzqxuk-59qg24sornk9qz73r1
o05fzsl 192.168.0.22:2377
Error response from daemon: remote CA does not match
fingerprint. Expected: 7867d80199d81783d39dfbaf321dd
72c380a3b6cf89785aad12f3d3ef3375c
(node2) (local) root@192.168.0.23 ~
$ docker swarm join --token SWMTKN-1-2y8p24lbdffaagks
w04jrjfoxm211nxuwiomye5wgl7fzqxuk-59qg24sornk9qz73r1
o05fzsl 192.168.0.22:2377
Error response from daemon: rpc error: code = Invalid
Argument desc = A valid join token is necessary to jo
in this cluster
(node2) (local) root@192.168.0.23 ~
$ docker swarm join --token SWMTKN-1-2y8p24lbdffaagks
w04jrjfoxm211nxuwiomye5wgl7fzqxuk-59qg24sornk9qz7er1
o05fzsl 192.168.0.22:2377
This node joined a swarm as a worker.
```

```
if the commandline doesn't appear in the terminal, make sure popups are enabled or try resizing the browser window.
```

```
h34cjef8a2m4zdrk8mvo35shr
(node1) (local) root@192.168.0.22 ~
$ docker network ls
NETWORK ID          NAME          DRIVER
SCOPE
3b2a510e9b6e       bridge       bridge
local
e7881e36d9b1       docker_gwbridge bridge
local
8d8bb63b82d2       host         host
local
nggzlbp002         ingress      overlay
swarm
7f2831962dc1       none         null
local
h34cjef8a2m4       overnet      overlay
swarm
(node1) (local) root@192.168.0.22 ~
$
w04jrjfoxm211nxuwiomye5wgl7fzqxuk-59qg24sornk9qz7er1
o05fzsl 192.168.0.22:2377
This node joined a swarm as a worker.
(node2) (local) root@192.168.0.23 ~
$ docker network ls
NETWORK ID          NAME          DRIVER
SCOPE
194a8f4b7418       bridge       bridge
local
74e1ae426a         docker_gwbridge bridge
local
4f0a4de27328       host         host
local
nggzlbp002         ingress      overlay
swarm
f09eb0ed70ae       none         null
local
(node2) (local) root@192.168.0.23 ~
$
```



```
$ docker network inspect overnet
[
  {
    "Name": "overnet",
    "Id": "h34cjef6a2m4zdrk8mvo35xhr",
    "Created": "2018-12-03T22:03:15.216402693Z",
    "Scope": "swarm",
    "Driver": "overlay",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "10.0.0.0/24",
          "Gateway": "10.0.0.1"
        }
      ]
    }
  }
]
```

Step 3: Create a service

Now that we have a Swarm initialized and an overlay network, it's time to create a service that uses the network.

Execute the following command from the first terminal to create a new service called *myservice* on the *overnet* network with two tasks/replicas.

```
docker service create --name myservice \
--network overnet \
--replicas 2 \
ubuntu sleep infinity
```

```
ov30itv6t2n7axy2goqbfqt5e
```

Verify that the service is created and both replicas are up by running `docker service ls`.

```
docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE
ov30itv6t2n7	myservice	replicated	2/2	ubuntu:latest

We can also run `docker network inspect overnet` on the second terminal to get more detailed information about the "overnet" network and obtain the IP address of the task running on the second terminal.

```
docker network inspect overnet
```

```
[
  {
    "Name": "overnet",
    "Id": "wlgmvajmmzskn84bqbd1tyuy",
    "Created": "2017-04-04T09:35:47.526642642Z",
    "Scope": "swarm",
    "Driver": "overlay",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "10.0.0.0/24",
          "Gateway": "10.0.0.1"
        }
      ]
    }
  }
]
```

```
(node1) (local) root@192.168.0.22 ~
$ docker service create --name myservice \
> --network overnet \
> --replicas 2 \
> ubuntu sleep infinity
i9pv2owh2qgsd88nra0fxb84
overall progress: 2 out of 2 tasks
1/2: running
2/2: running
verify: Waiting 3 seconds to verify that tasks are stable..
verify: Service converged
(node1) (local) root@192.168.0.22 ~
$ docker service ls
ID            NAME          MODE
REPLICAS     IMAGE         PORTS
i9pv2owh2qgs myservice     replicated
2/2          ubuntu:latest

(node1) (local) root@192.168.0.22 ~
$ !
```

```
$ docker network inspect overnet
[
  {
    "Name": "overnet",
    "Id": "h34cjef6a2m4zdrk8mvo35xhr",
    "Created": "2018-12-03T22:12:43.301222615Z",
    "Scope": "swarm",
    "Driver": "overlay",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "10.0.0.0/24",
          "Gateway": "10.0.0.1"
        }
      ]
    }
  }
],
```

```
"Internal": false,
"Attachable": false,
"Containers": {
  "fbc8bb0834429a68b2cccf25d3c90135dbda41e08b300f07845cb7f082bcd0f1": {
    "Name": "myservice.1.nicggj5tutar7h7sgsvqg72r",
    "EndpointID": "8edf83ebce77aed6d0193295c80c6aa7a5b76a08880a166002ecd3a2099bb6c",
    "MacAddress": "02:42:0a:00:00:03",
    "IPv4Address": "10.0.0.3/24",
    "IPv6Address": ""
  }
},
"Options": {
  "com.docker.network.driver.overlay.vxlanid_list": "4097"
}
```

```
"Containers": {
  "7e5e8653d867d324121fde541fd478e992650f652c8b0f4847a369096c57ed9b": {
    "Name": "myservice.1.jv114322r6flyjgdmme4shkj",
    "EndpointID": "efa732ef97670346861bae2bf3661230b9816d8b008216855f4f3691f19e7a54",
    "MacAddress": "02:42:0a:00:00:05",
    "IPv4Address": "10.0.0.5/24",
    "IPv6Address": ""
  }
},
```

Run a `docker ps` command to get the ID of the service task so that you can log in to it in the next step.

```
docker ps
```

CONTAINER ID	IMAGE	NAMES	COMMAND	CREATED
d676496d187	ubuntu@sha256:dd7808d8792c9841d0b460122f1acf0a2dd1f56404f8d1e56298048885e45535	myservice.1	"sleep infinity"	13 minutes ago
418e03fb77f8	ubuntu@sha256:dd7808d8792c9841d0b460122f1acf0a2dd1f56404f8d1e56298048885e45535	myservice.2	"sleep infinity"	Up 13 minutes

Log on to the service task. Be sure to use the container **ID** from your environment as it will be different from the example shown below. We can do this by running `docker exec -it <CONTAINER ID> /bin/bash`.

```
docker exec -it yourcontainerid /bin/bash
root@d676496d187:/#
```

Install the ping command and ping the service task running on the second node where it had a IP address of **10.0.3** from the `docker network inspect overnet` command.

```
[node1] (local) root@192.168.0.22 ~
$ docker ps
CONTAINER ID   IMAGE      COMMAND
CREATED       STATUS     PORTS
NAMES
418e03fb77f8   ubuntu:latest   "sleep infinity"
13 minutes ago Up 13 minutes
myservice.2.nnin2uulyeyl8ahl5tqq0f33a
[node1] (local) root@192.168.0.22 ~
$ docker exec -it 418e03fb77f8 /bin/bash
root@418e03fb77f8:/# apt-get update && apt-get install -y i
putils-ping
Get:1 http://security.ubuntu.com/ubuntu bionic-security InR
elease [83.2 kB]
Get:2 http://archive.ubuntu.com/ubuntu bionic InRelease [24
2 kB]
Get:3 http://security.ubuntu.com/ubuntu bionic-security/uni
```

Tutaj takie dwie wersje, bo nie wiem, o który terminal im chodziło. Pingowanie zarówno 10.0.0.5 jak i 10.0.0.6 z wnętrza przechodzi jak na drugim obrazku, choć wskazują roota wewnątrz. Z zewnątrz ping wygląda jak w rozwiązaniu, na obrazku pierwszym.

Now, lets ping **10.0.0.3**.

```
root@d676496d187:/# ping -c5 10.0.0.3
PING 10.0.0.3 (10.0.0.3) 56(84) bytes of data.
^C
--- 10.0.0.3 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 2998ms
```

The output above shows that both tasks from the **myservice** service are on the same overlay network spanning

```
[node2] (local) root@192.168.0.23 ~
$ ping -c5 10.0.0.5
PING 10.0.0.5 (10.0.0.5): 56 data bytes

--- 10.0.0.5 ping statistics ---
5 packets transmitted, 0 packets received, 100% packet loss
[node2] (local) root@192.168.0.23 ~
$
```

Install the ping command and ping the service task running on the second node where it had a IP address of **10.0.3** from the `docker network inspect overnet` command.

```
apt-get update && apt-get install -y iputils-ping
```

Now, lets ping **10.0.0.3**.

```
root@d676496d187:/# ping -c5 10.0.0.3
PING 10.0.0.3 (10.0.0.3) 56(84) bytes of data.
^C
--- 10.0.0.3 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 2998ms
```

The output above shows that both tasks from the **myservice** service are on the same overlay network spanning both nodes and that they can use this network to communicate.

Step 5: Test service discovery

Now that you have a working service using an overlay network, let's test service discovery.

If you are not still inside of the container log back into it with the `docker exec -it <CONTAINER ID> /bin/bash`

```
iputils-ping is already the newest version (3.20161105-1ubu
ntu2).
0 upgraded, 0 newly installed, 0 to remove and 4 not upgrad
ed.
root@418e03fb77f8:/# ping -c5 10.0.0.5
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=0.278 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=0.108 ms
64 bytes from 10.0.0.5: icmp_seq=3 ttl=64 time=0.150 ms
64 bytes from 10.0.0.5: icmp_seq=4 ttl=64 time=0.164 ms
64 bytes from 10.0.0.5: icmp_seq=5 ttl=64 time=0.179 ms

--- 10.0.0.5 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 399
7ms
rtt min/avg/max/mdev = 0.108/0.175/0.278/0.058 ms
root@418e03fb77f8:/# docker ps
bash: docker: command not found
root@418e03fb77f8:/#
```

Step 5

Try and ping the "myservice" name from within the container by running `ping -c5 myservice`.

```
root@d676496d187:/# ping -c5 myservice
PING myservice (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.020 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.052 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.044 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.042 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.056 ms

--- myservice ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4001ms
rtt min/avg/max/mdev = 0.020/0.042/0.056/0.015 ms
```

The output clearly shows that the container can ping the `myservice` service by name. Notice that the IP address returned is `10.0.0.2`. In the next few steps we'll verify that this address is the virtual IP (VIP) assigned to the `myservice` service.

Type the `exit` command to leave the `exec` container session and return to the shell prompt of your Docker host.

```
root@d676496d187:/# exit
```

```
root@418e03fb77f8:/# options ndots:0
bash: options: command not found
root@418e03fb77f8:/# ping -c5 myservice
PING myservice (10.0.0.4) 56(84) bytes of data.
64 bytes from 10.0.0.4 (10.0.0.4): icmp_seq=1 ttl=64 time=0.178 ms
64 bytes from 10.0.0.4 (10.0.0.4): icmp_seq=2 ttl=64 time=0.066 ms
64 bytes from 10.0.0.4 (10.0.0.4): icmp_seq=3 ttl=64 time=0.074 ms
64 bytes from 10.0.0.4 (10.0.0.4): icmp_seq=4 ttl=64 time=0.058 ms
64 bytes from 10.0.0.4 (10.0.0.4): icmp_seq=5 ttl=64 time=0.063 ms

--- myservice ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 399.8ms
rtt min/avg/max/mdev = 0.058/0.087/0.178/0.046 ms
```

Inspect the configuration of the "myservice" service by running `docker service inspect myservice`. Lets verify that the VIP value matches the value returned by the previous `ping -c5 myservice` command.

```
docker service inspect myservice
```

```
[
  {
    "ID": "ov30itv6t2n7axy2goqbft5e",
    "Version": {
      "Index": 19
    },
    "CreatedAt": "2017-04-04T09:35:47.009730798Z",
    "UpdatedAt": "2017-04-04T09:35:47.05475096Z",
    "Spec": {
      "Name": "myservice",
      "TaskTemplate": {
        "ContainerSpec": {
          "Image": "ubuntu:latest@sha256:dd7808d8792c9841d0b460122f1acf0a2dd1f56404f8d1e56298048885e45535",
          "Args": [
            "sleep",
            "infinity"
          ],
          "Endpoint": {
            "Spec": {
              "Mode": "vip"
            },
            "VirtualIPs": [
              {
                "NetworkID": "wlgqvajmmzskn84bqbd1ytuy",
                "Addr": "10.0.0.2/24"
              }
            ]
          }
        },
        "Resources": {}
      }
    },
    "Endpoint": {
      "Spec": {
        "Mode": "vip"
      },
      "VirtualIPs": [
        {
          "NetworkID": "wlgqvajmmzskn84bqbd1ytuy",
          "Addr": "10.0.0.2/24"
        }
      ]
    }
  }
]
```

Towards the bottom of the output you will see the VIP of the service listed. The VIP in the output above is `10.0.0.2` but the value may be different in your setup. The important point to note is that the VIP listed here matches the value returned by the `ping -c5 myservice` command.

Feel free to create a new `docker exec` session to the service task (container) running on `node2` and perform the same `ping -c5 service` command. You will get a response from the same VIP.

```
If the commandline doesn't appear in the terminal, make sure popup
"Mode": "vip"
}
},
"Endpoint": {
  "Spec": {
    "Mode": "vip"
  },
  "VirtualIPs": [
    {
      "NetworkID": "h34cjef6a2m4zdrk8mvo35xhr",
      "Addr": "10.0.0.4/24"
    }
  ]
}
]
(node1) (local) root@192.168.0.22 ~
$

5 packets transmitted, 0 packets received, 100% packet loss
(node2) (local) root@192.168.0.23 ~
$ docker network ls
NETWORK ID        NAME                DRIVER
SCOPE
194a8f4b7418      bridge             bridge
local
f74e1ae42f6a      docker_gwbridge    bridge
local
4f0a4de27328      host               host
local
nqqgzlpbp002      ingress            overlay
swarm
f09eb0ed70ae      none               null
local
h34cjef6a2m4      overnet            overlay
swarm
(node2) (local) root@192.168.0.23 ~
```

Cleaning Up

Feel free to create a new `docker exec` session to the service task (container) running on **node2** and perform the same `ping -c5 service` command. You will get a response from the same VIP.

Cleaning Up

Hopefully you were able to learn a little about how Docker Networking works during this lab. Lets clean up the service we created, the containers we started, and finally disable Swarm mode.

Execute the `docker service rm myservice` command to remove the service called *myservice*.

```
docker service rm myservice
```

Execute the `docker ps` command to get a list of running containers.

```
docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
NAMES					
846af8479944	ubuntu	"sleep infinity"	17 minutes ago	Up 17 minutes	h
euristic_boyd					
4e0da45b0f16	nginx	"nginx -g 'daemon ...'"	12 minutes ago	Up 12 minutes	443/tcp, 0.0.0.0:8080
->80/tcp	web1				

You can use the `docker kill <CONTAINER ID ...>` command to kill the ubuntu and nginx containers we started at the beginning.

```
docker kill yourcontainerid1 yourcontainerid2
```

Finally, lets remove node1 and node2 from the Swarm. We can use the `docker swarm leave --force` command to do that.

Lets run `docker swarm leave --force` on node1.

```
docker swarm leave --force
```

Lets also run `docker swarm leave --force` on node2.

```
docker swarm leave --force
```

Congratulations! You've completed this lab!

Share this on → [Twitter](#) [Facebook](#) [Google+](#) [LinkedIn](#)

```
If the commandline doesn't appear in the terminal, make
myservice.2.nnin2uulyeyl8ahl5tqq0f33a
[node1] (local) root@192.168.0.22 ~
$ docker kill 418e03fb77f8
Error response from daemon: Cannot kill container: 418e03fb
77f8: No such container: 418e03fb77f8
[node1] (local) root@192.168.0.22 ~
$ docker kill 418e03fb77f8
Error response from daemon: Cannot kill container: 418e03fb
77f8: No such container: 418e03fb77f8
[node1] (local) root@192.168.0.22 ~
$ docker ps
CONTAINER ID   IMAGE     COMMAND
CREATED        STATUS    PORTS
NAMES
[node1] (local) root@192.168.0.22 ~
$ docker swarm leave --force
Node left the swarm.
[node1] (local) root@192.168.0.22 ~
$

NETWORK ID     NAME      DRIVER
SCOPE
194a8f4b7418   bridge   bridge
local
f74e1ae42f6a   docker_gwbridge   bridge
local
4f0a4de27328   host     host
local
nqggzlpbp002   ingress  overlay
swarm
f09eb0ed70ae   none     null
local
h34cjef6a2m4   overnet  overlay
swarm
[node2] (local) root@192.168.0.23 ~
$ docker swarm leave --force
Node left the swarm.
[node2] (local) root@192.168.0.23 ~
$
```